N–Body Simulations of Barred Galaxies

———————————————————

A Thesis
Presented to
The Division of Mathematics and Natural Sciences
Reed College

———————————————————

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts

———————————————————

Thomas B Malthouse

Summer 2017

Approved for the Division
(Physics)

_____

J Powell

Acknowledgements

I want to thank a few people.

Preface

This is an example of a thesis setup to use the reed thesis document class.

List of Abbreviations

You can always change the way your abbreviations are formatted. Play around with it yourself, use tables, or come to CUS if you'd like to change the way it looks. You can also completely remove this chapter if you have no need for a list of abbreviations. Here is an example of what this could look like:

ABC	American Broadcasting Company
CBS	Columbia Broadcasting System
CDC	Center for Disease Control
CIA	Central Intelligence Agency
CLBR	Center for Life Beyond Reed
CUS	Computer User Services
FBI	Federal Bureau of Investigation
NBC	National Broadcasting Corporation

Table of Contents

List of Tables

List of Figures

Abstract

The preface pretty much says it all.

Dedication

You can have a dedication here if you wish.

Introduction

The Milky Way is an entirely unextraordinary galaxy. Its hallmark spiral arms—visible in the night sky as a bright smearing of stars, stretching from horizon to horizon—are shared by about 60% of galaxies in our universe (Loveday, 1996). A bar—a large collection of stars passing through the galactic center, prominent in renderings of the Milky Way (such as Fig. 1)—is also found around two in three other spiral galaxies. Under the Hubble classification system commonly used to sort and organize galaxies, the Milky Way is classified as an *Sb* type galaxy, along with 40% of known galaxies (for more information, see A.1). These extensive similarities mean that studying the evolution and structure of the Milky Way can provide insights about galactic behavior in general, and that observing other galaxies can reveal the past and future of out own.

Disks

The Milky Way has two main disks which hold the vast majority of visible matter in the galaxy and give rise to its spiral structure. The *thin disk* is the more visible of the two, composed mainly of main-sequence stars and clouds of gas and dust (Sparke & Gallagher, 2000). Its vertical density scale height—the distance over which its density decreases by a factor of e—is about $350\,\mathrm{pc}$—very thin compared to its radius of about $25\,\mathrm{kpc}$. This thinness comes from its young age, since the stars that compose it are less likely to have had their orbits perturbed—especially in the chaotic period about $9\,\mathrm{Gyr}$ ago. The thin disk accounts for about 97% of the galaxy's (normal) mass and holds nearly all galactic dynamism and stellar formation.

The other disk, referred to as the *dark disk* or *thick disk*, is far older and less dynamic (Moyer et al., 2003). Composed of stars formed $10\,\mathrm{Gyr}$ to $12\,\mathrm{Gyr}$ ago, it is very faint and hard to detect—all the bright stars burned out long ago, and the only ones left are low-magnitude red dwarfs and K-class stars. These stars' orbits also tend to be less regular, since they've had time to be perturbed and pushed into new orbits, especially during the chaotic initial organization of the Milky Way $10\,\mathrm{Gyr}$ ago—resulting in a scale height of about $1\,\mathrm{kpc}$. Because its stars are so steady-burning and the complete lack of gas and dust, the thick disk is very stable and exhibits none of the dynamism

Figure 1: This rendering of the Milky Way clearly shows the stellar bar going through the center of the galaxy, and the extensive spiral system. *Image courtesy NASA/JPL*

seen in the thin disk. It accounts for only about 3% of the regular matter in the galaxy, with a total mass of about $1 \times 10^{10}\,\mathrm{M}_\odot$. Note that the name *dark disk* refers to the low luminosity of the stars within, and has no relation to the dark matter discussed below.

Metallicity and Age

The easiest way to determine which disk a star is in is to look at its metallicity (the amount of metal in the star.) This can be measured by looking at the strength of various emission spectra, since elements like iron have very distinct emission lines. These heavy elements are only formed when large stars reach the end of their life, and so their concentration has steadily increased over time as more large stars form and die. Old stars dating back to the formation of the galaxy (like those of the thick disk) tend to have very "clean" emission spectra, with very little other than hydrogen and helium, while younger stars have strong magnesium and iron lines.

Metallicity isn't a perfect way to measure the age of a star. Metal concentrations vary widely across both space and time, and two stars forming at the same time may have very different metallicities. However, when looking at a large and statistically representative sample of stars, a high metallicity indicates a younger age (Carroll & Ostlie, 2006, pp. 885).

The Stellar Halo

The disks extend to about $25\,\mathrm{kpc}$ from the galactic center, and contain practically all the mass within that radius. Past that point, however, stars are distributed far more chaotically. The thin plane disappears, and stellar orbits become more spherically distributed. The composition of individual stars in the halo is similar to those in the thick disk—old, dim, and unchanging. However, the halo is also home to many globular clusters, groups of tens or hundreds of thousands of stars that act like small galaxies in their own right. These clusters continue to create new stars, and most of the light coming from the halo is from globular clusters.

Measurements of stellar velocities have long predicted that the mass of the halo dominates the mass of the galaxy—about 95% of galactic mass must be in the halo for the observed velocity curves to hold. Since the halo was known not to be made up of gas and dust (otherwise its extinctive properties would be easy to measure), astronomers long though that the halo contained vast numbers of dense, dark bodies, such as lone planets, dim stars, black holes, and neutron stars—referred to as MACHOs, or Massive Astrophysical Compact Halo Object. Gravitational lensing observations disproved

this theory, however, when they capped the mass percentage of MACHOs at about 15% of the mass of the galaxy. The remainder of the mass was some strange material, spherically distributed throughout the universe, that interacts with nothing but gravity.

Dark Matter

As that strange material was further studied (as much as it is possible to study something so elusive), more of its properties were discovered. This *dark matter* seems to be made up of WIMPS (Weakly Interacting Massive Particles), which only interact via gravity and the weak force. Being collisionless (since it does not interact with the electromagnetic force), it does not coalesce and form clouds and stars like regular matter, which is how it has maintained its spherical distribution for so long.

This dark matter has a density of

$$\rho(r) = \frac{\rho_0}{\left(\frac{r}{a}\right)\left(1 + \frac{r}{a}\right)^2} \tag{1}$$

where ρ_0 is the maximum density and a is proportional to the size of the galaxy. This function is referred to as the NFW profile (named after its creators, Navarro, Frank, and White), and is highly accurate for all observed spiral galaxies. For the Milky Way, the simplified equation

$$\rho(r) = \frac{\rho_0}{1 + \left(\frac{r}{1}\right)^2} \tag{2}$$

also returns satisfactory results. However, both these equations appear to suffer from a major problem. If we use them to calculate the total mass of the dark matter in a galaxy, integrating from 0 to ∞, it appears that a galaxy has an infinite mass, as follows:

$$\int\limits_{0}^{\infty} \rho(r) 4\pi r^2 \, \mathrm{d}r = \infty \tag{3}$$

Since we know this is not true, there must be a cutoff point where the law no longer holds. As it turns out, in local groups like our own, the dark matter halos are so large that they border one another—providing a natural cutoff point and a solution to our problem (Sparke & Gallagher, 2000, pp. 196).

Coordinate Systems

To identify and keep track of objects in the sky, we need to create a coordinate system. A number of natural possibilities spring to mind. The three most useful are detailed below.

Equatorial Coordinates

An equatorial coordinate has two components—a *right ascension*, and a *declination*. To find this coordinate, find the location on the Earth's surface where the object of interest is directly overhead, in the very middle of the sky. The right ascension is the angle between the nearest point on the equator and the vernal equinox (which is defined to be the point where the equator crosses the ecliptic). The declination is then the angle between the current point and that nearest equatorial point. This coordinate system is very ancient, dating back millennia. Figure 2 shows the process of finding such a coordinate.

Galactic Coordinates

This coordinate system is similar to the equatorial system, but the declination is measured from the galactic plane instead of the equatorial plane. The right ascension is then defined to be the angle between the projection of the body of interest on the galactic plane and the vector between the Sun and galactic center. Figure 3 shows this process in more detail. Although harder to calculate from the surface of the earth, this system is more natural when studying bodies traveling close to the Sun. The standard notation and conversions between the two systems are given below:

Measurement	Equatorial Notation	Galactic Notation
Right Ascension	δ	b
Declination	α	ℓ

$$\sin b = \sin \delta_{NGP} \sin \delta + \cos \delta_{NGP} \cos \delta \cos(\alpha - \alpha_{NGP}) \tag{4}$$

$$\sin \delta = \sin \delta_{NGP} \sin b + \cos \delta_{NGP} \cos b \cos \ell_{NCP} - \ell \tag{5}$$

Where $\delta_{NGP} = 27°7'41.7''$ and $\ell_{NCP} = 123°55'55.2''$, as determined by the tilt of the earth and its orientation relative to the galactic plane. These equations can also be inverted to find ℓ and α (Carroll & Ostlie, 2006, pp. 900).

Cylindrical Coordinates

The two coordinate systems discussed earlier are well-suited for positional observations from the Earth at a given point in time, but perform poorly over long timeframes. As the sun travels around its orbit, the coordinates of an object change even if it has not moved at all—not an ideal behavior from a reference system. The cylindrical coordinate system, with a reference point at the center of the universe solves these concerns. Unlike

Figure 2: This figure shows the process of finding an equatorial coordinate. CNP and CSP refer to the celestial north and south poles, respectively. R. Ascension refers to the right ascending node, where the Earth's rotational plane crosses its orbital plane.


Figure 3: This figure shows the process of finding a galactic coordinate. GNP and GSP refer to the galactic north and south poles, respectively.

the others, it is a three-component coordinate system: R is the radial distance along the plane, increasing outwards; θ is the angular position, and increases in the direction of rotation; and z is the height above (or below) the plane, increasing towards the north, as shown in figure 4. These coordinates also produce a natural (and commonly used) velocity coordinate system, as described below (Carroll & Ostlie, 2006):

$$\Pi \equiv \frac{dR}{dt} \qquad\qquad \Theta \equiv R\frac{d\theta}{dt} \qquad\qquad Z \equiv \frac{dz}{dt} \qquad (6)$$

Note that, because the galaxy rotates clockwise when viewed from the north pole, this is a left-handed coordinate system rather than a more-standard right-hand system. Fortunately, we do not need to take any cross-products, so this does not cause any problems.

Local Standard of Rest

Now that we have a definition of the cylindrical velocity coordinates, we can define the Local Standard of Rest (LSR), an important concept in astrophysics. The LSR at a given moment is defined to be the velocity of a body in the sun's position, in a perfectly circular and on-plane orbit—which in practice means the Θ-component of the sun's velocity, with Π and Z set to zero.

The velocity of a nearby star relative to the LSR is known as its *peculiar velocity*, and approximates the velocity of that star relative to the sun. Its coordinates are typically designated (u, v, w), where

$$u = \Pi - \Pi_{LSR} \qquad (7)$$
$$v = \Theta - \Theta_{LSR} \qquad (8)$$
$$w = Z - Z_{LSR} \qquad (9)$$

The average peculiar velocity for stars in the solar neighborhood is approximately zero, since the universe is mostly axisymmetric. However, individual peculiar velocities vary widely, with young main-sequence stars like the sun having low velocities and old, metal-poor red dwarfs having higher velocities. As discussed earlier, this is due to the additional orbital perturbations experienced by old stars, especially during the chaotic period of formation 9 Gyr ago.

Figure 4: This figures shows the process of finding a cylindrical coordinate.

Chapter 1

Computational Optimizations [[Perhaps a more introductory title – reserving this title for another section or chapter where is more appropriate]]

[[I am really glad to see this section start to develop, but when you get a chance to add a bunch of meta-discourse/introduction that would be great. Why C instead of C++ – type of discussion. Where learn C – "... the hard way." Some overall comments on programming and hardware for N-body simulations. Vectors are clearly very cool, but it seems a bit deep for a start. Indicate level of computer science education that you are writing for and give your old texts are refs. For example, compiler definition.]]

Simulating a galaxy requires a massive amount of computational power [[OK, even a neophyte has to know that there are a huge number of stars in a galaxy.]]. Even a simple statistical approximation, with a couple hundred or thousand bodies, [["statistical approximation" meaning? You are referring to the idea of approximating, say, 100, stars by one body, no?]] involves hundreds of thousands of calculations every timestep [[a lot going on in this sentence: first the approximation of numbers of stars, secondly the nature of the calculation – gravitation and $F = m\,a$ calculations? I looked this point is the first in the thesis where N–body calculations are mentioned, so start more basically or possibly ref. Noah]]. Multiplied by 30 updates [["updates" – a technical term define, because in the beginning that same word could be interpreted by a tyro in its generic sense]] per second (to ensure smooth graphics), we quickly become responsible for hundreds of millions of calculations every second [[nice point – good use of "responsible"]]. To achieve that kind of performance, we need to understand how modern computer

hardware functions and how various programming languages take advantage of that [[lots of writers don't use referents, but it's fairly safe to say that referents – especially in a thesis – are an excellent habit.]].

1.0.1   Assembly and Machine Code

[[Perhaps a basic intro. couple of sentences laying out the hierarchy of computer programs. Include that thorough and pedagogical definition of those two kinds of code?]]

Every computer program ever written can be expressed [[rewritten ? out]] with a small set of [[need a good modifier here to foreshadow machine code?]] commands. These [[very basic commands]] commands directly manipulate memory, adding, comparing, and copying bits around [[at the level of the computer that is just barely above individual transistors]]. Different types of computer processors have different sets of commands, some more extensive than others [[some sentence close to this for that intro. sentence]]. For example, [[this example is pretty good – althought the square root reference is a real zinger ...]] the [[always define an acronym – as you did with RISC below]] ARM processors found in smartphones have a *RISC* (reduced instruction set circuit) design, which means they can perform a minimal set of operations, and more complex operations—like the square root—are expressed in terms of that small set. Other architectures, like the x86 design found in nearly all modern laptops and desktops, are referred to as *CISC* (complex instruction set circuit) designs, and can do more operations without being programmed to. This design is often faster, but is more difficult to manufacture and often more expensive. In this thesis, we will exclusively focus on the x86_64 architecture, as found in nearly all modern desktop and laptop computers [[nice sentence – maybe an introductory-partner sentence way up at the top.]].

The set of operations a processor [[very good stuff, but it still need a little bit of support – even a reference to Simpson's chapter, but you might be able to do it better yourself in a couple of sentences by paraphrasing Simpsons]] can perform is referred to as the *instruction set*, and is determined entirely by the physical circuit of the processor—so an x86 processor can never be turned into an ARM processor, or vice versa. As the processor runs, instructions are provided in the form of 8-bit numbers, each one corresponding to an operation. For example, the *opcode* `0x40` (decimal 64) tells the processor to increment a given value by one. Since the processor also needs to know what data to perform the operation on, a more complete instruction would look something like `0x40 6c`. This instruction increments the value stored in memory at location 6c by one.

Early programs were written [[with combinations of instructions]] like this [[one in the paragraph above]], as long sequences of hexadecimal numbers that could be di-

rectly interpreted by the processor [["directly interpreted by the processor" I believe that I know what that clause means, but support in for the uninitiated.]]. However, remembering dozens of opcodes was unwieldy, and finding and fixing bugs was practically impossible [[interesting claim – why exactly?]]. Frustrated, programmers invented human-friendly names for operations and common memory locations [[good, but needs support]] and wrote a tool (called an *assembler*) to automatically replace them with the appropriate number [[which appropriate number? the one that refers to a location in RAM]]. This new programming language was called *assembly* [[You can see the reader's issue the definition comes after a bunch of tough stuff.]], and greatly simplified programming. For example, our instruction from [[the]] earlier [[example of a cell phone]] would be written in assembly as something like `INC EAX` —still opaque, but more useful than a few hexadecimal numbers. Many programs still in use today (including Unix, the basis of Mac OS X and Linux) were originally written in an assembly language [[cool]].

1.0.2   The Development of High–level Languages

Assembly was a massive improvement over typing in raw machine code, but it remained difficult to write and maintain large codebases [[define – seems intuitive]], especially with large teams of programmers. Simple tools like conditional (if/else) blocks and loops were error-prone [[seems as though you must have gotten this knowledge from some reference, no? If so include it here.]], and the structure of the program mimicked the inner working of the processor more than that of the human brain [[structure of the human brain ... dangerous territory ...]]. Various *programming languages* were designed to simplify the programming process, with associated *compilers* to convert source code into assembly [[seems good]]. These languages featured more natural math and now–familiar constructs like the `for` loop and the concept of functions [[very familiar from MATLAB – at least for me, define/explain]]. The most popular and enduring of these languages was C, which underlies nearly every operating system and utility [[define pedagogically]] in use today. [[Foreshadow the fundamental differences between C and, say, Python or MATLAB]] Our example instruction from earlier, written in C, is `i++;` —far simpler than before, and close to the notation used in mathematics.

Of course (and this [[a well-constructed referent]] should seem familiar by this point [[not quite understanding ... point of the thesis?]]), C had some limitations. String [[five-word Sci. Am. definition]] handling was an exercise in frustration, and manually handling memory meant leaks (where a piece of memory is never returned to the operating system [[good definition]]) were commonplace. Accessing invalid memory was all-too-common [[surely a reference would be welcome here by someone really trying to understand the claim at a deep level]], especially among new programmers, and could lead to all kinds of insidious bugs. Higher-level languages were created to deal with these problems, providing robust, error-free constructs to simplify these tasks.

The first and most well-known of these languages was C++, which started out as a collection of text replacements and code generators for C and quickly evolved into a massive, full-featured language. Others followed, many keeping the basic syntax of C but automating ever-more. In general, these languages were much easier to write than C, at a significant performance cost [[classic claim. explain why – harks back to the difference between complied and not-compiled languages, no?]] —which was offset by rapid improvements in processor technology [[at least until recently, i.e. the last couple of years.]].

1.0.3   Choosing a Language

As mentioned earlier, performance is absolutely critical for our n–body simulator. Therefore, we need to be very careful when choosing a programming language for our project. Modern *Interpreted languages*—which do not compile to machine code, but are run by an *interpreter* or *virtual machine*—like Python or C# have too much overhead [[true, but explain]] to use [[this excellent paragraph could have been foreshadowed, thereby helping the reader.]]. Their simplicity and ease of use would be very nice to have, but they just aren't fast enough for our purposes. That leaves us with the compiled languages—C and C++ the most well known among them [[at least mention fortran, no?]].

At first, C++ would seem like an obvious choice. Features like exceptions (which greatly improve error handling) and inheritance (which makes it easier to work with complex data types) are hard to live without, and most programmers have never used a language without them. However, these features also add overhead—they function as code generators, making it easy to create a lengthy and computationally complex program without realizing it. Since we can't use most of the features that make C++ what it is for performance reasons, the choice between C and C++ is mostly cosmetic— the finished codes would look very similar, and could easily be converted into each other. Since I had more experience with C from prior project and much prefer the C syntax to that of C++, I decided to use C for this project.

1.0.4   Learning C

C has a bit of a reputation. It's seen as a fragile, temperamental language, where even the smallest change can make the difference between a crash and a successful run. Pointers—the defining language feature—can be difficult to understand at first, and are far less user-friendly than the objects and arrays most languages replace them with. But it's also a very simple language, where learning a dozen keywords and a handful of concepts lets you understand just about any C program ever written. There are countless resources online for learning C, and it isn't hard to start writing simple programs.

For a [[annotated]] list of useful resources, see A.2.

1.1   Hardware-based Vectors

Let's imagine the components of an n-body simulator. We clearly need to define some kind of data type to hold essential information about a given body—its position, its velocity (collectively referred to as its *state*), its mass, and possible some kind of name or numeric ID, to allow individual bodies to be tracked. In C, we would define this `Body` type using a `struct`, or compound data type. The latter fields are simple enough—C provides datatypes for both integral and non-integral numbers, and strings are easy to create. However, defining the vectors (as needed for the position and velocity) is a little trickier. A simple definition for `Body` is shown in listing 1.1.

> Listing 1.1 :   This code snippet shows the definition of the `Body` data type. Note that the position and velocity are of the `vec3` datatype, which is not part of the C language.

```c
//We create a new datatype, and declare it as a compound type
typedef struct {
  //The first element is an unsigned 32-bit integer,
  //which holds a unique ID for the body
  uint32_t id;
  //We then have a 3-dimensional vector to hold the body's position
  vec3 pos;
  //And another to hold its velocity
  vec3 vel;
  //Then, we have its mass, stored as a floating-point number.
  double mass;
//Finally, we give the compound structure the name 'Body'
} Body;
```

In the past, to write a physical simulation involving vector calculations, a programmer would have to create their own definition of a vector using a `struct`, like we did to define a `Body` (a sample definition is shown in 1.2). They would be responsible for coding even the most basic operations—like addition and scalar multiplication—by hand, and this code would be executed like any other, with the normal overhead of a function call. However, recent x86-64 processors (anything made by Intel in the last decade) have a single-instruction, multiple-data (SIMD) coprocessor (Intel, 2016), which is designed to perform a given operation on multiple numbers simultaneously.

> Listing 1.2 :   An old-style definition of a vector

```c
//Like before, we create a new compound data structure
typedef struct {
    //With 3 floating-point elements, named x, y, and z
```

```
    double x;
    double y;
    double z;
//And name the whole structure 'vec3'
} vec3;
```

1.1.1   SIMD

SIMD, short for *single-instruction, multiple-data*, is the simplest form of parallelism. A special piece of hardware executes the same command on more than one piece of data at a time, making it very useful for tasks like array processing or graphics rendering, where each element in the array or each pixel on the screen needs the same computations performed on it. For tasks like array processing, the SIMD optimizations are often generated by the compiler, with no programmer input necessary (other than setting the correct flags). This model—where each element is computed simultaneously—would also seem to work well for basic vector operations like addition and scalar multiplication, which are calculated on an element-by-element basis. Since the compiler cannot figure out whether a given data type is meant to represent a vector, we need to explicitly tell it. The syntax for doing so is shown in listing 1.3.

Listing 1.3 : The declaration for EXT vector types. These vectors are not part of standard C, but are available on nearly all modern systems using the `clang` compiler.

```
    //This declares a new datatype called vec3.
    //It consists of double-precision floating point numbers,
    //and is defined to be a 3-dimensional vector by the
    //__attribute__ tag.
    typedef double vec3 __attribute__((ext_vector_type(3)));
```

1.1.2   Properties of Hardware-based Vectors

Although these types are meant to serve as vectors, many useful vector operations—absolute value, the cross and dot products, and similar—are not implemented. As a general rule, any operation that returns a scalar (like the absolute value and dot product, for example), or doesn't have a scalar analogue (like the cross product, which is only defined on 3-dimensional vectors), must be implemented by the programmer. Writing the code for these operations tends to be fairly simple, and uses the standard algorithms covered in Physics 101. For example, listing 1.4 shows the code for calculating the absolute value of a vector.

Listing 1.4 : Implementing the absolute value function in software

```
    inline double vabs (vec3 v) {
```

```
        return sqrt(v.x*v.x + v.y*v.y + v.z*v.z);
   }
```

Some operations are already predefined. Since these vectors follow the SIMD model, any operation that acts on each element individually is trivial for the compiler to handle. Therefore, all of the normal C operators are valid for vectors. Sometimes, like in the case of addition or scalar multiplication, this behavior is normal and expected, and provides a useful speed improvement over a software-based method. However, some operators make no sense under this model yet are defined anyways. For example, a vector space is not an ordered field (as we learned in linear algebra). There is no good way to order or sort vectors. And yet the `< >` operators are defined for vectors, behaving as shown in listing 1.5. This behavior can be a little tricky, especially when an operator behaves unexpectedly.

> Listing 1.5 : Some operations aren't especially useful

```
(vec3){0,5,8} > (vec3){3,4,8} == (vec3){false, true, false};
```

1.1.3   Performance Advantages

Now that we have our two ways of defining vectors, we need to compare their performance. The best way to measure the performance of various kinds of vectors is to write a simple program and time it, both with and without hardware vectors enabled. However, looking at the assembly output of the compiler can also provide valuable information, since each instruction has a similar execution time. For example, listing 1.6 shows the code for a performance critical loop that calculates force, compiled with the option `-Ofast` —which enables maximum optimizations, but still uses the old–style vectors from earlier.

Listing 1.7, on the other hand, has been compiled with vector support—using options `-Ofast -march=native` telling the compiler to take advantage of any and all CPU features that could be useful. It is about two-thirds as long as the fastest possible software version, which generally corresponds to a 33% performance increase—a huge boon in code that runs tens of millions of times per second. Rows and rows of near-identical `add` and `mov` calls are replaced by a single specialized call to `vfmadd231sd` , a vector-based instruction that performs all those operations simultaneously.

[[Need a lot of help with these listing –> if not they go to the appendix]]

> Listing 1.6 :   The optimized assembly for a critical loop, without SIMD vector support. Don't worry if you can't understand this—neither can I. The important thing is its relative length.

```
LBB2_27:                                    ##    in Loop: Header=BB2_23 Depth=1
```

```
        movapd  144(%rsi), %xmm2
        movapd  160(%rsi), %xmm3
        movsd   128(%rsi), %xmm4        ## xmm4 = mem[0],zero
        subpd   48(%rsp), %xmm3         ## 16-byte Folded Reload
        subpd   64(%rsp), %xmm2         ## 16-byte Folded Reload
        movapd  %xmm2, %xmm5
        mulsd   %xmm5, %xmm5
        movapd  %xmm2, %xmm6
        shufpd  $1, %xmm6, %xmm6        ## xmm6 = xmm6[1,0]
        mulsd   %xmm6, %xmm6
        movapd  %xmm3, %xmm7
        mulsd   %xmm7, %xmm7
        addsd   %xmm5, %xmm7
        addsd   %xmm6, %xmm7
        xorps   %xmm5, %xmm5
        sqrtsd  %xmm7, %xmm5
        addsd   %xmm10, %xmm7
        divsd   %xmm7, %xmm4
        movddup %xmm4, %xmm6            ## xmm6 = xmm4[0,0]
        movddup %xmm5, %xmm7            ## xmm7 = xmm5[0,0]
        divsd   %xmm5, %xmm3
        divpd   %xmm7, %xmm2
        mulpd   %xmm6, %xmm2
        mulpd   %xmm4, %xmm3
LBB2_28:                               ##   in Loop: Header=BB2_23 Depth=1
        addpd   %xmm3, %xmm0
        addpd   %xmm2, %xmm1
```

Listing 1.7 : The same code, with SIMD vectors

```
LBB2_25:                               ##   in Loop: Header=BB2_21 Depth=1
        vmovupd 160(%rsi), %ymm2
        vmovsd  144(%rsi), %xmm3        ## xmm3 = mem[0],zero
        vsubpd  (%rsp), %ymm2, %ymm2    ## 32-byte Folded Reload
        vpermilpd       $1, %xmm2, %xmm4 ## xmm4 = xmm2[1,0]
        vmulsd  %xmm4, %xmm4, %xmm4
        vfmadd231sd     %xmm2, %xmm2, %xmm4
        vextractf128    $1, %ymm2, %xmm5
        vfmadd231sd     %xmm5, %xmm5, %xmm4
        vaddsd  %xmm0, %xmm4, %xmm6
        vdivsd  %xmm6, %xmm3, %xmm3
        vbroadcastsd    %xmm3, %ymm3
        vsqrtsd %xmm4, %xmm4, %xmm4
        vmovddup        %xmm4, %xmm6    ## xmm6 = xmm4[0,0]
        vdivpd  %xmm6, %xmm2, %xmm2
        vdivsd  %xmm4, %xmm5, %xmm4
        vinsertf128     $1, %xmm4, %ymm2, %ymm2
        vfmadd231pd     %ymm2, %ymm3, %ymm8
```

1.1.4 Disadvantages

There is a good reason SIMD support is not enabled by default. Because its instructions are not part of the `x86-64` standard, various chipset manufacturers may implement the feature differently (or not enable it at all.) The `-march=native` flag we passed to the compiler voids any guarantee that the resulting binary be able to run on any computer running the same operating system and instruction set. We could build the non-SIMD code in 32-bit mode and load the resulting binary onto a computer from the early 90s, and it would (probably) run. Trying the same with the SIMD-enabled binary may throw an error similar to the one seen in 1.8, especially on old or lower-end CPUs.

> Listing 1.8 : An error thrown by an unsupported instruction on OS X. Running vector code on an unsupported processor would lead to something like this.

```
45584 illegal hardware instruction  ./a.out
```

1.2   The Tree Algorithm

By implementing hardware-based vectors, we've sped up our force-calculating function by nearly 30%. However, running that function still takes *some* amount of time, and that will remain true no matter how much we optimize it. We can—and should—fiddle around with low-level assembly all day, but at some point that strategy faces diminishing returns. There are only so many optimizations that can be squeezed out of a piece of code, and at some point we will reach the shortest possible runtime for an algorithm. At that point, the only way to speed up our computation is to do less of it, replacing chunks of calculations with good-enough approximations that take a fraction of the time to calculate.

When trying to calculate the behavior of a body in a field, the simplest way is to calculate the force on that body from every other body that makes up that field. This strategy is easy to implement—the only relevant equations are from Physics 101. However, as the number of bodies to be considered grows, the required number of computations skyrockets, becoming untenable on even the most powerful computers. In technical terms, this algorithm is said to have a *time efficiency* of $O(n^2)$, which means that, as the number of bodies doubles, the number of calculations quadruples.

We clearly need to find a way to reduce the number of calculations. Using the system in figure 1.1 as an example, we see that each body requires 79 calculations—a good baseline to work from. Recall from Physics 101 that, as one moves further away, the forces exerted by two identical bodies near each other converge. It follows that, if distant collections of bodies could be replaced with singular large bodies, the number of calculations could be sharply reduced without significantly impacting accuracy. Figure

1.2 shows an example of this, replacing a cluster of 10 distant bodies and reducing the number of calculations by about 15%—with only one replacement.

Our next task is to find an algorithm to do all these merges and accuracy judgements as quickly as possible. Representing space as a kind of tree structure makes this easy to do. To build this tree, space is divided into nodes along each axis (so a 2-dimensional node will have 4 children, while a 3-dimensional one will have 8). Each of these child nodes is then bisected in the same manner, and the process repeats recursively until each bottom-level node has 0 or 1 bodies in it. The resulting structure is referred to as a quadtree (in 2D space) or an octotree (in 3D space). The first two levels of a quadtree can be seen in figure 1.3.

Then, to create a system that approximates the original, we simply have to "walk" through the tree, deciding whether each node is a suitably accurate approximation for the bodies contained within, or whether its children should be considered individually. This decision is based on the mass of the node and its distance from the body of interest, as well as an "accuracy factor" that controls the allowable level of error. Eventually, a list of nodes is created that contains each body exactly once, and the net force can be found by summing up the forces from each node. Figure

The algorithm discussed here has a time complexity of $O(n \log n)$ in the average case, and $O(n^2)$ in the very worst case. This means that it will practically always beat the naïve algorithm for systems of more than a few bodies, and becomes very valuable for large simulations. The creation and analysis of the tree does add some overhead, but a well written implementation will rarely take longer than a few dozen force calculations—while eliminating hundreds or even thousands.

Figure 1.1: We want to calculate the force on the red dot in this system, which involves 79 calculations (since there are 80 bodies). These 79 calculations must be repeated for every other body, for a total of $79 * 80 = 6320$ calculations.


Figure 1.2: The force on the red body in this system is similar to the one seen earlier. However, the number of calculations has been significantly reduced, because about 10 bodies have been replaced with one large body.

Figure 1.3: The quadtree after two levels of bisections. Note the 4 child nodes of each node, and the ease of scaling this structure to an arbitrary dimension.


Figure 1.4: This system is a good approximation of the one considered earlier, and only requires 24 calculations—nearly 75% less than before. Of course, a real algorithm would be far more rigorous, and likely find even more opportunities to merge bodies.

Conclusion

The process by which the structure and dynamics of the MW were discovered was by no means trivial: Linblad [ref.] was a hero along with the other pioneers. Buried in the galaxy they had some advantages compared to understanding, say, the Andromeda, but being in the disk causes huge difficulties, not the least of which is the dust.

4.1 More info

And here's some other random info: the first paragraph after a chapter title or section head *shouldn't be* indented, because indents are to tell the reader that you're starting a new paragraph. Since that's obvious after a chapter or section title, proper typesetting doesn't add an indent there.

Appendices

Appendix A

A.1   The Hubble Classification System

The Hubble Classification System, or Hubble Sequence, is the most commonly used scheme for classifying galaxies. Nicknamed *The Fork*, it goes from elliptical galaxies on the left, to the two kinds (barred and unbarred) of spiral galaxies on the right.


Image courtesy Allan Sandage/CalTech

A.2   C Language Resources

- *Learn C the Hard Way: Practical Exercises on the Computational Subjects You Keep Avoiding (Like C)* —A good introduction to writing simple, useful C programs. Often breaks with tradition, but always for the better.

  `http://www.indiebound.org/book/9780321884923`

- *C in a Nutshell* —A massive, 824 page book, that has everything you could ever want to know about using the C language

  `http://www.indiebound.org/book/9781491904756`

- *21st Century C* —A guide to using a language designed in 1970, in 2017. Most useful if you already have an understanding of C.

  `http://www.indiebound.org/book/9781491903896`

- *The C Programming Language* —An absolute classic, written by the designers of the language. This book isn't all that useful for *learning* modern C, but it sheds light on the reasons the language was designed the way it was.

  `http://www.indiebound.org/book/9780131103627`

- *Build your own Lisp* —Learn C by writing an interpreter for the Lisp language. A complete guide to a really fun and useful project. Available both online and as a printed book.

  `http://www.buildyourownlisp.com`

  `http://www.indiebound.org/book/9781501006623`

References

Carroll, B. W., & Ostlie, D. A. (2006). *An Introduction to Modern Astrophysics (2nd Edition)*. Pearson. `https://www.amazon.com/Introduction-Modern-Astrophysics-2nd/dp/0805304029?SubscriptionId=0JYN1NVW651KCA56C102&tag=techkie-20&linkCode=xm2&camp=2025&creative=165953&creativeASIN=0805304029`

Debattista, V. P., Ness, M., Gonzalez, O. A., Freeman, K., Zoccali, M., & Minniti, D. (2016). Separation of stellar populations by an evolving bar: Implications for the bulge of the milky way. ArXiv:1611.09023.

Erwin, P., & Debattista, V. P. (2016). Caught in the act: Direct detection of galactic bars in the buckling phase. *The Astrophysical Journal Letters*, *825*(2), L30. `http://stacks.iop.org/2041-8205/825/i=2/a=L30`

Eskridge, P. B., Frogel, J. A., Pogge, R. W., Quillen, A. C., Davies, R. L., DePoy, D. L., Houdashelt, M. L., Kuchinski, L. E., RamÃŋrez, S. V., Sellgren, K., Terndrup, D. M., & Tiede, G. P. (2000). The frequency of barred spiral galaxies in the near-infrared. *The Astronomical Journal*, *119*(2), 536. `http://stacks.iop.org/1538-3881/119/i=2/a=536`

Gardner, E., Debattista, V. P., Robin, A. C., Vásquez, S., & Zoccali, M. (2014). N-body simulation insights into the x-shaped bulge of the milky way: kinematics and distance to the galactic centre. *Monthly Notices of the Royal Astronomical Society*, *438*(4), 3275. `+http://dx.doi.org/10.1093/mnras/stt2430`

Intel (2016). Intel streaming simd extensions technology defined. `https://www.intel.com/content/www/us/en/support/processors/000005779.html`

Loveday, J. (1996). The apm bright galaxy catalogue. *Monthly Notices of the Royal Astronomical Society*, *278*(4), 1025. `+http://dx.doi.org/10.1093/mnras/278.4.1025`

Malthouse, T. (????). Trajectory: an n-body simulator in c, using simd vector extensions and opencl for maximum performance. `https://github.com/wisdomgroup/trajectory`.

Merritt, D., & Sellwood, J. A. (1994). Bending instabilities in stellar systems. *Astrophysical Journal*, *425*, 551–567.

Moyer, E., Sion, E. M., Szkody, P., GÃďnsicke, B., Howell, S., & Starrfield, S. (2003). Hubble space telescope observations of the old nova di lacertae. *The Astronomical Journal*, *125*(1), 288. http://stacks.iop.org/1538-3881/125/i=1/a=288

of Washington, U. (????). ChaNGa (Charm N-body GrAvity solver). https://github.com/N-BodyShop/changa/wiki/ChaNGa.

Sparke, L. S., & Gallagher, J. S. (2000). *Galaxies in the Universe: An Introduction*. Cambridge University Press. https://www.amazon.com/Galaxies-Universe-Introduction-Linda-Sparke/dp/0521592410?SubscriptionId=0JYN1NVW651KCA56C102&tag=techkie-20&linkCode=xm2&camp=2025&creative=165953&creativeASIN=0521592410